


Generative Adversarial Networks

From Kullback-Leibler Divergence to Vanilla Generative Adversarial Networks


List of contents

- Kullback-Liebr Divergence
- Jensen-Shannon Divergence
- Generative Adversarial Networks (Introduction)
- Generative Adversarial Networks (Formulation)
- Generative Adversarial Networks (Properties)

Acknowledgements: [Lilian Weng](#) from OpenAI & [François Chollet](#) from Google



What is a Generative Adversarial Network?



Kullback-Liebert Divergence

Kullback-Liebert Divergence (KLD) aims to measure how a probability distribution p diverges from a second expected probability distribution q so that:

$$D_{KL}(p||q) = \int_x p(x) \log \frac{p(x)}{q(x)} dx$$

Kullback-Liebert Divergence


Kullback-Liebert Divergence (KLD) aims to measure how a probability distribution p diverges from a second expected probability distribution q so that:

$$D_{KL}(p||q) = \int_x p(x) \log \frac{p(x)}{q(x)} dx$$

D_{KL} achieves the minimum when $p(x) = q(x)$


Kullback-Lieber Divergence

According to its formulation, KLD is asymmetric by design



Kullback-Lieber Divergence

According to its formulation, KLD is asymmetric by design



In cases where $p(x)$ is close to zero while $q(x)$ is significantly non-zero, the $q(x)$ effect is ignored. This property may cause undesirable results when we just want to measure the similarity between two equally important distributions

Jensen-Shannon Divergence

Jensen-Shannon Divergence (JSD) aims to measure how a probability distribution p diverges from a second expected probability distribution q so that:

$$D_{JS}(p\|q) = \frac{1}{2} D_{KL} \left(p \parallel \frac{p+q}{2} \right) + \frac{1}{2} D_{KL} \left(q \parallel \frac{p+q}{2} \right)$$

Jensen-Shannon Divergence


Jensen-Shannon Divergence (JSD) aims to measure how a probability distribution p diverges from a second expected probability distribution q so that:

$$D_{JS}(p\|q) = \frac{1}{2} D_{KL} \left(p \parallel \frac{p+q}{2} \right) + \frac{1}{2} D_{KL} \left(q \parallel \frac{p+q}{2} \right)$$

D_{JS} achieves the minimum when $p(x) = q(x)$

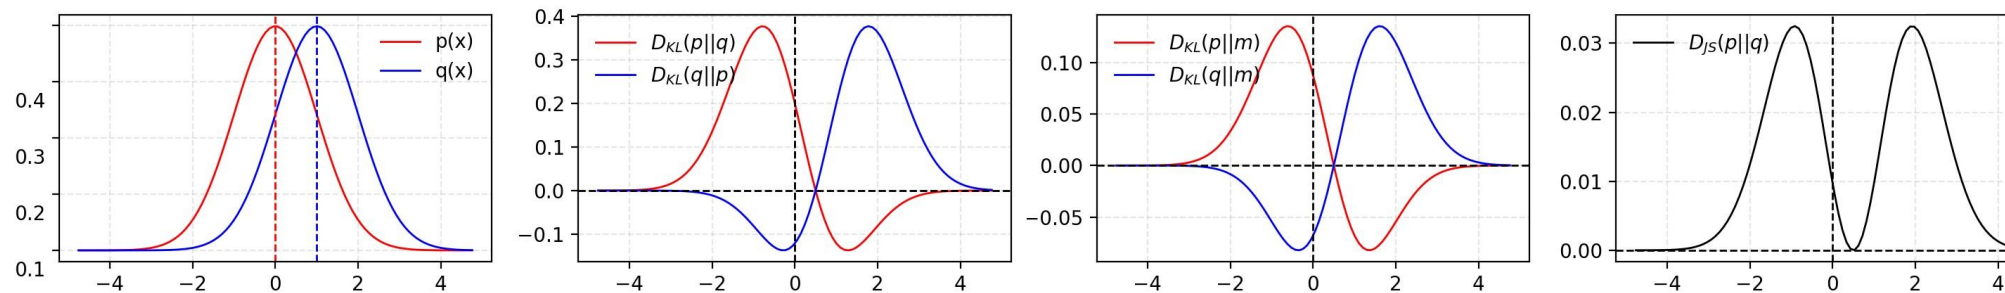
Jensen-Shannon Divergence

According to its formulation, JSD is symmetric by design, is bounded in the $[0,1]$ range, and is smoother than KLD



Jensen-Shannon Divergence

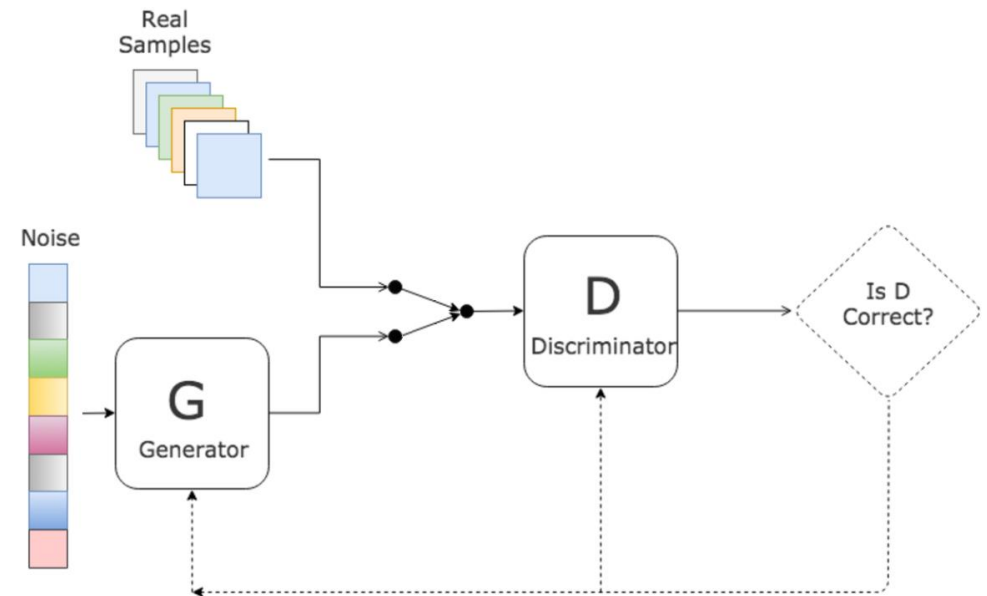
According to its formulation, JSD is symmetric by design, is bounded in the $[0,1]$ range, and is smoother than KLD



One reason behind Generative Adversarial Networks success is switching the loss function from asymmetric KLD in traditional maximum-likelihood approach to symmetric JSD

Generative Adversarial Networks (Introduction)


A Generative Adversarial Network (GAN) is an artificial neural network composed of two models:



Generative Adversarial Networks (Introduction)

A Generative Adversarial Network (GAN) is an artificial neural network composed of two models:

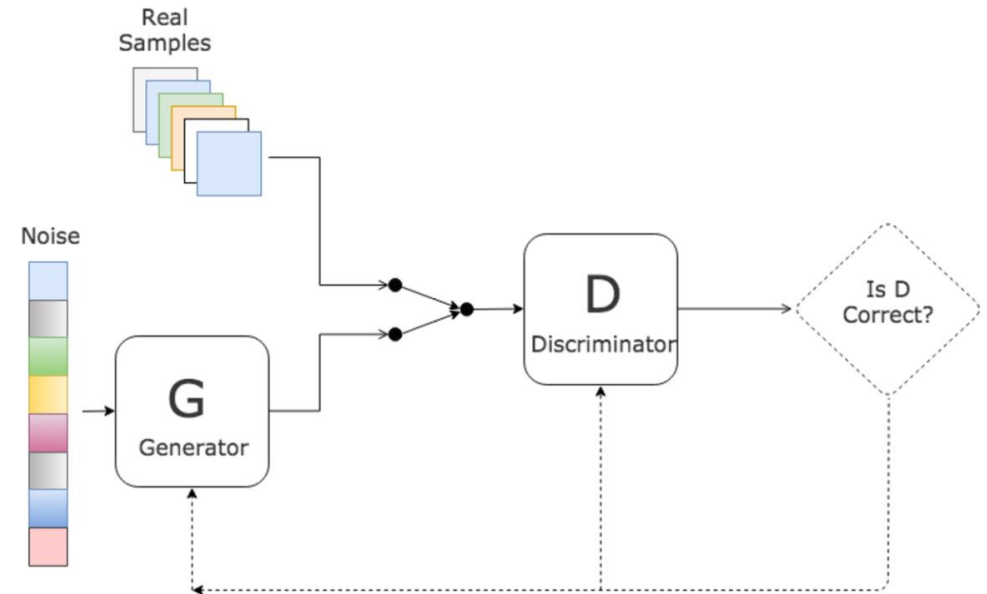
- A discriminator D to estimate the probability of a given sample coming from the real dataset. It works as a critic and is optimized to discriminate the fake samples from the real ones



Generative Adversarial Networks (Introduction)

A Generative Adversarial Network (GAN) is an artificial neural network composed of two models:

- A discriminator D to estimate the probability of a given sample coming from the real dataset. It works as a critic and is optimized to discriminate the fake samples from the real ones
- A generator G to output synthetic samples given a noise variable input z . It is trained to capture the real data distribution so that its generative samples can be as real as possible, or in other words, can trick the discriminator to offer a high probability



Generative Adversarial Networks (Formulation)

Given a noise variable z (usually normally or uniformly distributed) and a dataset x , we define the following distributions:

Generative Adversarial Networks (Formulation)

Given a noise variable z (usually normally or uniformly distributed) and a dataset x , we define the following distributions:

- p_z as the data distribution over noise z

Generative Adversarial Networks (Formulation)

Given a noise variable z (usually normally or uniformly distributed) and a dataset x , we define the following distributions:

- p_z as the data distribution over noise z
- p_g as the generator's distribution over data x

Generative Adversarial Networks (Formulation)

Given a noise variable z (usually normally or uniformly distributed) and a dataset x , we define the following distributions:

- p_z as the data distribution over noise z
- p_g as the generator's distribution over data x
- p_r as the data distribution over real samples x

Generative Adversarial Networks (Formulation)

We want to improve the discriminator D ability to accurately recognize real data coming from distribution $p_r(x)$ maximizing $\mathbb{E}_{x \sim p_r(x)} [\log D(x)]$

Generative Adversarial Networks (Formulation)

We want to improve the discriminator D ability to accurately recognize real data coming from distribution $p_r(x)$ maximizing $\mathbb{E}_{x \sim p_r(x)} [\log D(x)]$

Moreover, given the fake samples $G(z) \sim p_z(z)$ is expected to output a probability $D(G(z))$ close to zero maximizing $\mathbb{E}_{z \sim p_z(z)} \left[\log (1 - D(G(z))) \right]$

Generative Adversarial Networks (Formulation)

We want to improve the discriminator D ability to accurately recognize real data coming from distribution $p_r(x)$ maximizing $\mathbb{E}_{x \sim p_r(x)} [\log D(x)]$

Moreover, given the fake samples $G(z) \sim p_z(z)$ is expected to output a probability $D(G(z))$ close to zero maximizing $\mathbb{E}_{z \sim p_z(z)} \left[\log \left(1 - D(G(z)) \right) \right]$

The generator is instead trained to increase its chances of producing good quality fake examples and with a high probability managing to deceive D . For this purpose, its goal is to minimize $\mathbb{E}_{z \sim p_z(z)} \left[\log \left(1 - D(G(z)) \right) \right] = \mathbb{E}_{x \sim p_g(x)} [\log(1 - D(x))]$

Generative Adversarial Networks (Formulation)

D and G are playing a minimax game in which we should optimize the following loss function:

$$\min_G \max_D L(G, D) = \mathbb{E}_{x \sim p_r(x)} [\log(D(x))] + \mathbb{E}_{x \sim p_g(x)} [\log(1 - D(x))]$$

Generative Adversarial Networks (Formulation)

D and G are playing a minimax game in which we should optimize the following loss function:

$$\min_G \max_D L(G, D) = \mathbb{E}_{x \sim p_r(x)} [\log(D(x))] + \mathbb{E}_{x \sim p_g(x)} [\log(1 - D(x))]$$

where $\mathbb{E}_{x \sim p_r(x)} [\log(D(x))]$ has no impact on G during gradient descent updates

Generative Adversarial Networks (Properties)

Given the minmax loss function just defined, the optimal value of $D(x)$ can be found solving the following integral:

$$L(G, D) = \int_x \left(p_r(x) \log(D(x)) + p_g(x) \log(1 - D(x)) \right) dx$$

Generative Adversarial Networks (Properties)

Given the minmax loss function just defined, the optimal value of $D(x)$ can be found solving the following integral:

$$L(G, D) = \int_x \left(p_r(x) \log(D(x)) + p_g(x) \log(1 - D(x)) \right) dx$$

Let $\tilde{x} = D(x)$, $A = p_r(x)$, $B = p_g(x)$. Since we can safely ignore the integral as x is sampled over all the possible values, the argument of the integral $f(\tilde{x})$ becomes:

$$f(\tilde{x}) = A \log(\tilde{x}) + B \log(1 - \tilde{x})$$

Generative Adversarial Networks (Properties)

By deriving with respect to \tilde{x} we obtain that:

$$\frac{df(\tilde{x})}{d\tilde{x}} = A \frac{1}{\ln 10} \frac{1}{\tilde{x}} - B \frac{1}{\ln 10} \frac{1}{1-\tilde{x}} = \frac{1}{\ln 10} \frac{A-(A+B)\tilde{x}}{\tilde{x}(1-\tilde{x})}$$

Generative Adversarial Networks (Properties)

By deriving with respect to \tilde{x} we obtain that:

$$\frac{df(\tilde{x})}{d\tilde{x}} = A \frac{1}{\ln 10} \frac{1}{\tilde{x}} - B \frac{1}{\ln 10} \frac{1}{1-\tilde{x}} = \frac{1}{\ln 10} \frac{A-(A+B)\tilde{x}}{\tilde{x}(1-\tilde{x})}$$

Thus, by setting $\frac{df(\tilde{x})}{d\tilde{x}} = 0$, we get the best value for the discriminator:

$$D^*(x) = \tilde{x} = \frac{A}{A+B} = \frac{p_r(x)}{p_r(x)+p_g(x)} \in [0,1]$$

Generative Adversarial Networks (Properties)

By deriving with respect to \tilde{x} we obtain that:

$$\frac{df(\tilde{x})}{d\tilde{x}} = A \frac{1}{\ln 10} \frac{1}{\tilde{x}} - B \frac{1}{\ln 10} \frac{1}{1-\tilde{x}} = \frac{1}{\ln 10} \frac{A-(A+B)\tilde{x}}{\tilde{x}(1-\tilde{x})}$$

Thus, by setting $\frac{df(\tilde{x})}{d\tilde{x}} = 0$, we get the best value for the discriminator:

$$D^*(x) = \tilde{x} = \frac{A}{A+B} = \frac{p_r(x)}{p_r(x)+p_g(x)} \in [0,1]$$

Once the generator is trained to its optimal $p_g \approx p_r$ and then $D^*(x) = \frac{1}{2}$

Generative Adversarial Networks (Properties)

When both G and D are at their optimal values, we have $p_g = p_r$ and $D^*(x) = \frac{1}{2}$ and the loss function becomes:

$$L(G^*, D^*) = \int_x \left(p_r(x) \log(D^*(x)) + p_g(x) \log(1 - D^*(x)) \right) dx$$

Generative Adversarial Networks (Properties)

When both G and D are at their optimal values, we have $p_g = p_r$ and $D^*(x) = \frac{1}{2}$ and the loss function becomes:

$$\begin{aligned} L(G^*, D^*) &= \int_x \left(p_r(x) \log(D^*(x)) + p_g(x) \log(1 - D^*(x)) \right) dx \\ &= \log \frac{1}{2} \int_x p_r(x) dx + \log \frac{1}{2} \int_x p_g(x) dx \end{aligned}$$

Generative Adversarial Networks (Properties)

When both G and D are at their optimal values, we have $p_g = p_r$ and $D^*(x) = \frac{1}{2}$ and the loss function becomes:

$$\begin{aligned} L(G^*, D^*) &= \int_x \left(p_r(x) \log(D^*(x)) + p_g(x) \log(1 - D^*(x)) \right) dx \\ &= \log \frac{1}{2} \int_x p_r(x) dx + \log \frac{1}{2} \int_x p_g(x) dx \\ &= -2 \log 2 \end{aligned}$$

Generative Adversarial Networks (Properties)

According to the previously defined formula, the Jensen-Shannon Divergence between p_r and p_g can be written as:

$$D_{JS}(p_r \| p_g) = \frac{1}{2} D_{KL} \left(p_r \| \frac{p_r + p_g}{2} \right) + \frac{1}{2} D_{KL} \left(p_g \| \frac{p_r + p_g}{2} \right)$$

Generative Adversarial Networks (Properties)

According to the previously defined formula, the Jensen-Shannon Divergence between p_r and p_g can be written as:

$$\begin{aligned} D_{JS}(p_r \| p_g) &= \frac{1}{2} D_{KL} \left(p_r \parallel \frac{p_r + p_g}{2} \right) + \frac{1}{2} D_{KL} \left(p_g \parallel \frac{p_r + p_g}{2} \right) \\ &= \frac{1}{2} \left(\log 2 + \int_x p_r(x) \log \frac{p_r(x)}{p_r(x) + p_g(x)} dx \right) + \frac{1}{2} \left(\log 2 + \int_x p_g(x) \log \frac{p_g(x)}{p_r(x) + p_g(x)} dx \right) \end{aligned}$$

Generative Adversarial Networks (Properties)

According to the previously defined formula, the Jensen-Shannon Divergence between p_r and p_g can be written as:

$$\begin{aligned} D_{JS}(p_r \| p_g) &= \frac{1}{2} D_{KL} \left(p_r \parallel \frac{p_r + p_g}{2} \right) + \frac{1}{2} D_{KL} \left(p_g \parallel \frac{p_r + p_g}{2} \right) \\ &= \frac{1}{2} \left(\log 2 + \int_x p_r(x) \log \frac{p_r(x)}{p_r(x) + p_g(x)} dx \right) + \frac{1}{2} \left(\log 2 + \int_x p_g(x) \log \frac{p_g(x)}{p_r(x) + p_g(x)} dx \right) \\ &= \frac{1}{2} (\log 4 + L(G, D^*)) \end{aligned}$$

Generative Adversarial Networks (Properties)

Essentially the loss function of a GAN quantifies the similarity between the generative data distribution p_g and the real sample distribution p_r by JSD when the discriminator is optimal:

$$L(G, D^*) = 2D_{JS}(p_r \| p_g) - 2\log 2$$

Generative Adversarial Networks (Properties)

Essentially the loss function of a GAN quantifies the similarity between the generative data distribution p_g and the real sample distribution p_r by JSD when the discriminator is optimal:

$$L(G, D^*) = 2D_{JS}(p_r \| p_g) - 2\log 2$$

The best G^* that replicates the real data distribution leads the loss function to reach the following minimum:

$$L(G^*, D^*) = -2\log 2$$